

Evaluating AI-Generated Occasion Outfits for Indian Contexts

Wedding, Festival & Office Wear | Josh Talks AI Product Task, Question 1

1. Objective and Use Case

This evaluation focuses on the Fashion / Apparel category, specifically occasion-specific outfit inspiration images for Indian contexts — the kind of feature a styling or shopping app would use to suggest "what to wear" for a wedding, a festival, or the office, rather than a product catalog photo of one specific item. The goal was to determine which AI image model correctly reflects Indian dressing norms for a given occasion instead of defaulting to a generic Western silhouette.

2. Why This Use Case Matters for India

India's dress codes are sharply occasion- and region-specific — a wedding guest outfit, a Diwali outfit, and an office outfit are governed by very different norms, and even "wedding guest" varies substantially by region (North vs. South Indian weddings, for instance). Outfit-inspiration is an increasingly common discovery feature in Indian shopping and styling apps, but it is only useful if the model understands these distinctions rather than returning a generic cocktail dress or Western business suit regardless of what was asked.

3. Why an AI Lab Building for India Should Care

Most global image models are trained predominantly on Western fashion imagery. This assessment surfaces three specific competency questions for the lab:

- Does the model default to a Western silhouette (a suit, a cocktail dress) even when explicitly asked for Indian occasion wear?
- Can the model distinguish between genuinely different Indian dress codes — festive-casual (Diwali at home) vs. formal-festive (wedding guest) vs. professional (office)?
- Does a regional cue ("South Indian wedding") change the output at all, or does it get flattened into a generic "Indian wedding" look?

4. How the Evaluation System Works

A custom, web-based blind rating tool was built and run with 10 participants.

- Blind testing: for each occasion, all model outputs were shown together labeled only "Option A / B / C" in randomized order, so raters never knew which model made which image.
- Likert scoring: each option was individually rated 1–5 on four criteria — occasion-appropriateness, cultural/context fit, styling coherence, and wearability.
- Forced choice: after scoring, raters answered "Which outfit would you actually pick?" — a single holistic pick per occasion, used to compute win rate.
- Coverage: 5 occasions (2 wedding-guest genders, Diwali, corporate office wear, and a regionally specific South Indian wedding) were each rated by 10 participants across all 3 required models — OpenAI GPT Image 1, Google Gemini 2.5 Flash Image, and Google Gemini 3.1 Flash-Lite Image — for 53 total outfit-level judgment instances (expected 50; see the data-quality note in Section 5).

5. Results





Figure 1 — Overall win rate across 53 outfit judgments, verified directly from the raw ratings log.

Overall leaderboard

Rank	Model	Win rate	Overall (1–5)	Wins
1	OpenAI — GPT Image 1	77%	3.67	41
2	Google — Gemini 3.1 Flash-Lite Image	13%	3.48	7
3	Google — Gemini 2.5 Flash Image	9%	3.47	5


Per-criterion averages



Model	Occasion-appropriateness	Cultural/context fit	Styling coherence	Wearability
GPT Image 1	3.64	3.53	3.81	3.70
Gemini 3.1 Flash-Lite Image	3.83	3.49	3.30	3.30
Gemini 2.5 Flash Image	3.75	3.60	3.21	3.32

GPT Image 1 leads on styling coherence and wearability by a clear margin (3.81 / 3.70) despite Gemini 3.1 Flash-Lite Image scoring marginally higher on raw occasion-appropriateness (3.83 vs. 3.64). This gap between individual-criterion scores and the forced-choice win rate (77% vs. 13%) is itself a finding: raters gave Gemini 3.1's outputs credit for matching the occasion in the abstract, but overwhelmingly preferred GPT Image 1's outputs when asked to actually pick one — suggesting GPT Image 1's advantage lives in overall styling execution and wearability rather than any single labeled criterion.

Wins by outfit



Outfit	GPT Image 1	Gemini 2.5 Flash Image	Gemini 3.1 Flash-Lite Image
Wedding Guest — Woman	8	2	0
Wedding Guest — Man	7	1	2
Diwali at Home	7	2	1
Corporate Office Wear — Woman	8	0	2
South Indian Wedding Guest	11	0	2

GPT Image 1 won every single outfit category outright, including the regionally specific South Indian Wedding Guest prompt — its strongest category by raw win count. This is a useful negative-space finding: it suggests none of the three models struggled specifically with the regional cue relative to the generic wedding-guest prompt (GPT Image 1's win count is actually highest there), so the earlier hypothesis that "South Indian" would be flattened into generic "Indian wedding" imagery was not clearly supported by this sample — though see the data-quality note below before reading too much into that category specifically.

Data-quality note

Four of the five outfit categories sum to exactly 10 judgments, matching the 10 registered participants. The South Indian Wedding Guest category sums to 13 (11 + 0 + 2), three more than expected. This is most likely a small number of duplicate submissions (e.g., a participant re-submitting after a page refresh) rather than a genuine 13th and 14th rater, since no additional names appear in the consent log beyond the 10 listed. The finding above is reported as-is, but should be treated as slightly noisier than the other four categories until the raw log is de-duplicated by participant email.

6. How to Scale This Further

The next stage would be to move from stated preference to revealed preference — running the winning outfit images as actual "outfit of the day" suggestions inside a styling app and measuring click-through and save/wishlist rate among real Indian users, rather than relying solely on a judgment panel. Other scaling directions:

- Expand the occasion bank to cover more regional and religious contexts (Bengali wedding, Eid, Onam, a North-East Indian festival) to test regional coverage more thoroughly than a single South Indian data point.
- Automate image generation via each model's API on a fixed occasion-prompt set so new model versions can be benchmarked on a recurring cadence.
- Recruit raters segmented by gender and region so preferences can be broken out by who is actually being asked to wear the outfit.
- De-duplicate submissions by participant email/device at the ingestion layer to avoid the kind of double-count seen in the South Indian Wedding Guest category above.

7. UI / Dashboard / Leaderboard

The evaluation tool itself doubles as the product prototype for this deliverable. It has three parts: an admin setup screen where per-model, per-occasion image URLs are pasted in; a blind rating flow where participants score each option on the four criteria and then make a single holistic pick after giving consent; and a live leaderboard showing win rate, per-criterion averages, and a wins-by-outfit breakdown, plus the participant/consent log for audit.

Final Reflection

- 1. Why did you choose this eval:** Outfit generation tests a well-documented, checkable failure mode — Western-default bias in occasion dressing — rather than a purely subjective "which looks nicer" question, and it maps directly onto a real product feature (styling/outfit discovery) that Indian shopping apps are actively building.
- 2. Why is it useful for India:** Indian dress codes are occasion- and region-specific in a way global fashion training data likely under-represents. An outfit-suggestion feature that keeps defaulting to Western silhouettes is not usable for an Indian audience, and this eval directly tests whether that default persists.
- 3. Why would an AI lab building for India care about it:** It isolates a specific, fixable gap — cultural/contextual fit in occasion dressing — separate from general image quality, giving the lab an actionable target rather than a vague "make it more Indian" note.
- 4. What did you learn from running the sample:** GPT Image 1 won decisively on forced choice (77%) despite a near-tie on individual Likert criteria with Gemini 3.1 Flash-Lite Image, showing that holistic "would you pick this" preference and componentized scoring can diverge — a model can score well on labeled dimensions while still losing the outfit someone would actually choose. The data-quality anomaly in the South Indian Wedding Guest category is also a concrete lesson in why de-duplication needs to be built into the ingestion pipeline, not just the analysis.
- 5. What would you improve with more time:** De-duplicate the ratings log by participant identity before drawing conclusions from any single-category result, and move beyond stated preference toward a revealed-preference test — surfacing the winning outfit images as real in-app suggestions and measuring save/click behavior among Indian users.